

ECE 449/590 – OOP and Machine Learning

Lecture 11 Tensor Class Design

Professor Jia Wang
Department of Electrical and Computer Engineering
Illinois Institute of Technology

September 28, 2022

Outline

Assertion and Exception

Tensor

Tensor Class Design

Reading Assignment

- ▶ This lecture: Accelerated C++ 9
- ▶ Next lecture: Accelerated C++ 13

Outline

Assertion and Exception

Tensor

Tensor Class Design

Enforcing Preconditions

- ▶ Class invariants help to enforce preconditions regarding data members for member functions.
 - ▶ What about preconditions regarding other function parameters?
- ▶ Use comments to document precondition around the function declaration.
 - ▶ Comments are in natural languages, which are usually ambiguous when the preconditions are very complicated.
 - ▶ Programmers may violate it accidentally even if they follow the instruction.
 - ▶ You have to update the comments as you update the function implementation.
- ▶ A function may be called many times during execution, some with correct arguments and some without.
 - ▶ Compiler/linker are not quite helpful in such cases (as of now).
 - ▶ Runtime validations are a must.

Assertion

- ▶ A C++ feature that allows the debugger to break your program when a precondition is violated.
 - ▶ The program simply prints an error message and exits if a debugger is not presented.
- ▶ The program will break at a point depending on the library implementations of assertions.
 - ▶ There may exist many implementations of assertions that provide different diagnosis informations.
- ▶ You can always use the call stack to navigate to your code that causes the violation before resolving it.
 - ▶ The assertion itself may provide useful information on what causes the violation.

The assert Macro

```
#include <assert.h>
date::date(int y, int m, int d):
    year_(y), month_(m), day_(d) {
    assert(valid());
}
```

- ▶ You can write your own assertions by using the `assert` macro from the standard header `assert.h` .
 - ▶ A macro is a piece of code like a function.
- ▶ `assert` takes an expression as the argument and will produce an informative message if the expression evaluates to false.
 - ▶ Assume `valid()` returns `true` if the object is valid.
 - ▶ If there is a debugger, it will also be triggered.

Exceptions

```
date::date(int y, int m, int d):  
    year_(y), month_(m), day_(d) {  
    if (!valid()) {  
        throw std::runtime_error("Invalid date");  
    }  
}
```

- ▶ The violation of the precondition can be notified by throwing an exception.
 - ▶ You can `throw` an object of any type.
 - ▶ Though in practice, people design different class types for different reasons of errors.
- ▶ `std::runtime_error` is a class type from the standard header `stdexcept` indicating an error at runtime.
 - ▶ Recall that `std::runtime_error("Invalid date")` constructs a `std::runtime_error` object with the argument `"Invalid date"`.
- ▶ There are other standard exception types.

Exception Handling

```
void some_function() {  
    date someday(2019, 2, 29);  
    ... // usual business flow  
}  
bool do_business() {  
    try {  
        // usual business flow  
        some_function();  
        another_function();  
    }  
    catch (std::exception &e) {  
        std::cerr << e.what() << std::endl;  
        return false;  
    }  
    return true;  
}
```

- ▶ Though you are forced to handle exceptions in your program, you don't need to handle them immediately.
 - ▶ Improve readability by NOT flooding usual business flows with error handling codes
 - ▶ If an exception is not handled, the program will abort.

Returned Error Codes vs Assertions vs Exceptions

- ▶ Returned error codes: function may also choose to return error codes to indicate violation of preconditions.
 - ▶ Can recover from an error.
 - ▶ No enforcement of error handling.
 - ▶ Awkward to return error codes from many functions.
 - ▶ Mixing code for business logic and error handling will affect readability.
- ▶ Assertions
 - ▶ No error recovery.
 - ▶ Enforce error handling by terminating program.
 - ▶ May be turned off to improve performance.
- ▶ Exceptions
 - ▶ Can recover from an error.
 - ▶ Enforce error handling by terminating program if not handled.
 - ▶ Prefer centralized error handling to improve readability.
- ▶ Choice between exceptions and assertions is a design decision. You need to make trade-offs.

Outline

Assertion and Exception

Tensor

Tensor Class Design

- ▶ Multidimensional data structure widely used for machine and deep learning.
 - ▶ E.g. a video can be represented as a tensor with 4 dimensions: frame, x, y, color
- ▶ Can be treated as generalization of scalars, vectors, and matrices.
 - ▶ Which are tensors of 0, 1, and 2 dimensions respectively.
- ▶ Performance of many machine and deep algorithms is closely related to efficiency of the underlying tensor operations.
 - ▶ Factors affecting efficiency include memory layout, cache architecture, parallel implementation, etc.
 - ▶ Let's focus on memory layout for this lecture.

Memory Layout

- ▶ The way to store a multidimensional tensor in memory.
- ▶ Need to consider many trade-offs.
 - ▶ Implementation of tensor operations with some kinds of layouts are more efficient than others.
 - ▶ Different libraries and languages may support different kinds of layouts.
 - ▶ Converting from one kind of layout to another usually requires to make copy of the data, consuming substantial amount of time even when there is enough memory.

(Contiguous) Row-Major Order

- ▶ A widely used memory layout for dense matrices.
- ▶ Consider a matrix A with R rows and C columns.
 - ▶ Let $A(i, j)$ be the elements on the i th row and j th column.
 - ▶ For simplicity we assume 0 based indices, i.e.
 $i = 0, 1, \dots, R - 1$ and $j = 0, 1, \dots, C - 1$.
- ▶ Store A in an array *data* row-by-row:

$$\begin{aligned} &A(0, 0), A(0, 1), \dots, A(0, C - 1), \\ &A(1, 0), A(1, 1), \dots, A(1, C - 1), \\ &\dots, \\ &A(R - 1, 0), A(R - 1, 1), \dots, A(R - 1, C - 1) \end{aligned}$$

- ▶ The array *data* has $R * C$ elements.
- ▶ $A(i, j) = data[i * C + j]$

Contiguous Row-Major Order for Tensors

- ▶ Dimension of the tensor A : N
- ▶ Shape of A : s_0, s_1, \dots, s_{N-1}
 - ▶ Can be stored in an array of N elements.
- ▶ One element: $A(i_0, i_1, \dots, i_{N-1})$
 - ▶ $i_k = 0, 1, \dots, s_k - 1$ for $k = 0, 1, \dots, N - 1$
- ▶ Store A in an array *data* as:

$A(0, \dots, 0, 0), A(0, \dots, 0, 1), \dots, A(0, \dots, 0, s_{N-1} - 1),$

$A(0, \dots, 1, 0), A(0, \dots, 1, 1), \dots, A(0, \dots, 1, s_{N-1} - 1),$

$\dots,$

$A(s_0 - 1, \dots, s_{N-2} - 1, 0), A(s_0 - 1, \dots, s_{N-2} - 1, 1), \dots, A(s_0 - 1, \dots, s_{N-2} - 1, s_{N-1} - 1)$

- ▶ The array *data* has $s_0 * s_1 * \dots * s_{N-1}$ elements.
- ▶ $A(i_0, i_1, \dots, i_{N-1}) = \text{data}[i_0 * s_1 * \dots * s_{N-1} + i_1 * s_2 * \dots * s_{N-1} + \dots + i_{N-1}]$.

Passing Tensors from Python to C

```
extern "C" int add_op_param_ndarray(  
    program *prog, const char *key,  
    int dim, size_t shape[], double data[]);
```

- ▶ Use contiguous row-major order.
 - ▶ `dim` is the dimension N of the tensor.
 - ▶ The `shape` array contains s_0, \dots, s_{N-1} .
 - ▶ The `data` array contains the elements of the tensor.
- ▶ Some implementation details
 - ▶ NumPy `ndarrays` are converted to such format when necessary.
 - ▶ As Python GC may release those buffers after this function return, you should make copies of `shape` and `data` arrays.

Passing Tensors from C to Python

```
extern "C" int execute(evaluation *eval,  
    int *p_dim, size_t **p_shape, double **p_data);
```

- ▶ We make use of pointer-to-pointers to allow Python code to access the shape and data arrays in C code.
- ▶ Since Python code will access those two arrays after this function returns, the two arrays need to have lifetimes beyond this function.
- ▶ The same function can return a scalar when necessary.
 - ▶ Recall scalars are tensors of dimension 0.
 - ▶ e.g. for Project 2

What about multidimensional arrays in C?

- ▶ There are other methods to support multidimensional arrays in C, e.g.
 - ▶ Built-in C multidimensional arrays.
 - ▶ Multiple levels of pointer-to-pointers.
- ▶ Not easy to work with
 - ▶ Need to specify dimension and/or shape at compile time.
 - ▶ Need to work with multiple levels of pointers and pointer arithmetics.
- ▶ Let's focus on contiguous row-major order and how to manage it in C++.

Outline

Assertion and Exception

Tensor

Tensor Class Design

Class Invariant

- ▶ To represent a tensor in contiguous row-major order, we need,
 - ▶ `dim`, `shape`, and `data`.
 - ▶ As data members of `tensor` class.
- ▶ Class invariant
 - ▶ `dim` is a positive integer.
 - ▶ The `shape` array should have `dim` elements, which are all positive integers.
 - ▶ The `data` array should have `shape[0]*shape[1]*...*shape[dim-1]` elements.
- ▶ What about scalars?

Updated Class Invariant

- ▶ If `dim == 0`, then a scalar is stored in `tensor`.
 - ▶ The `shape` array should be empty.
 - ▶ The `data` array should have a single element being the scalar.
- ▶ If `dim > 0`, then a tensor is stored in `tensor`.
 - ▶ The `shape` array should have `dim` elements, which are all positive integers.
 - ▶ The `data` array should have `shape[0]*shape[1]*...*shape[dim-1]` elements.

The tensor Class

```
class tensor {
public:
    tensor(); // scalar 0
    explicit tensor(double v); // scalar v
    tensor(int dim, size_t shape[], double data[]); // from C
    ...
private:
    std::vector<size_t> shape_;
    std::vector<double> data_;
}; // class tensor
```

- ▶ Use ctors to establish class invariant.
 - ▶ Use `std::vector` to store arrays.
 - ▶ There is no need to store `dim` as it can be obtained as `shape_.size()`.
- ▶ Ctors taking a single parameter should usually be declared as `explicit` to prevent implicit conversions that may cause hard-to-debug issues.

Implementing Ctors

```
tensor::tensor():  
    data_(1, 0) {  
}  
  
tensor::tensor(double v):  
    data_(1, v) {  
}  
  
tensor::tensor(int dim, size_t shape[], double data[]):  
    shape_(shape, shape+dim) {  
        // calculate N as shape[0]*shape[1]*...*shape[dim-1]  
        ...  
        data_.assign(data, data+N);  
}
```

- Similar to how we handle [inputs](#) for Project 2, C arrays can be copied into C++ vectors using ctors or [assign](#).

Accessors

```
class tensor {  
public:  
    ...  
    int get_dim() const;  
  
    // scalar only  
    double item() const;  
    double &item();  
    ...  
}; // class tensor
```

- ▶ Allow users of `tensor` to know the dimension by `get_dim`.
- ▶ Scalar can be accessed by `item` – note that two versions are provided.
 - ▶ The `const item` will be called with `const tensor` objects, only allowing to read the scalar.
 - ▶ The non-`const item` will be called with other `tensor` objects, allowing to read and write the scalar via the reference.

Implementing const and Non-const Members

```
double tensor::item() const {  
    assert(shape_.empty());  
    return data_[0];  
}
```

```
double &tensor::item() {  
    assert(shape_.empty());  
    return data_[0];  
}
```

- ▶ Although the two function bodies look exactly the same, they are actually different since different `[]` operators are used.
 - ▶ We will study how to implement our own `vector` later.
- ▶ Use assertions to make sure that the `tensor` object indeed holds a scalar.
 - ▶ You may choose to use exceptions here as well.
- ▶ What about accessors for tensors?

More Accessors

```
class tensor {
public:
    ...
    double at(size_t i) const;
    double at(size_t i, size_t j) const;
    ...
}; // class tensor

double tensor::at(size_t i) const {
    assert(get_dim() == 1);
    assert(i < shape_[0]);
    return data_[i];
}

double tensor::at(size_t i, size_t j) const {
    assert(get_dim() == 2);
    assert((i < shape_[0]) && (j < shape_[1]));
    return data_[i*shape_[1]+j];
}
```

- ▶ We may create accessors for specific dimensions.
 - ▶ Many tensor operations we need to implement have specific requirements on tensor dimensions.
- ▶ Use assertions or exceptions to guard against misuse.

Passing tensor Back to C Code

```
class tensor {
public:
    ...
    size_t *get_shape_array();
    double *get_data_array();
    ...
}; // class tensor

size_t *tensor::get_shape_array() {
    return shape_.empty()? nullptr: &shape_[0];
}
double *tensor::get_data_array() {
    return &data_[0];
}
```

- ▶ `std::vector` provides backward compatibility with C arrays.
- ▶ However, this feature should be used with care.

Managing tensor Lifetime

```
class evaluation {  
    ...  
    tensor &get_result();  
private:  
    ...  
    std::map<int, tensor> variables_;  
}; // class evaluation
```

- ▶ Intermediate variables as **tensors** are stored within **evaluation**.
- ▶ They will be there as long as the **evaluation** object is not destroyed and **variables_** is not cleared.
- ▶ **get_result** will return one of them corresponding to the result of the evaluation.

Updated `execute()` Function

```
int execute(evaluation *eval,
            int *p_dim, size_t **p_shape, double **p_data)
{
    ... // logging and error checking
    tensor &res = eval->get_result();
    *p_dim = res.get_dim();
    *p_shape = res.get_shape_array();
    *p_data = res.get_data_array();
    return 0;
}
```

- ▶ Please modify the code once you are done with Project 2.
- ▶ Our Python code will be able to construct a NumPy `ndarray` as result after `execute` returns.

Summary and Advice

- ▶ Use assertions and exceptions to enforce error handling.
- ▶ Tensors are multidimensional arrays.
- ▶ Tensors can be conveniently stored and passed around in contiguous row-major order.
- ▶ Always start your class design with a well designed class invariant.