

ECE 449/590 – OOP and Machine Learning

Lecture 25 OpenMP

Professor Jia Wang
Department of Electrical and Computer Engineering
Illinois Institute of Technology

November 21, 2022

Outline

OpenMP Overview

Hello World!

Single Program Multiple Data

Reading Assignment

- ▶ This lecture: Introduction to OpenMP
 - ▶ Most of today's lecture is based on Tim Mattson's presentation "The OpenMP Common Core: A hands on exploration".
 - ▶ The whole presentation and the slides can be found at <https://www.youtube.com/watch?v=I2EaVMjZRRY>
- ▶ Next week: no in-person lectures
 - ▶ Please watch the video "Introduction to GPU Architecture and Programming Models", Tim Warburton, Virginia Tech <https://www.youtube.com/watch?v=uvVy3CqpVbM>

Outline

OpenMP Overview

Hello World!

Single Program Multiple Data

Parallel and Distributed Computing

- ▶ Parallel and distributed computing are two different paradigms to utilize computing resources beyond a single core.
 - ▶ Multiple cores available from a modern processor.
 - ▶ Multiple machines connected by a network.
- ▶ Parallel computing: Single Program Multiple Data (SPMD)
 - ▶ All cores the same binary program, though each may work on a different chunk of data.
- ▶ Distributed computing: coordination of heterogeneous tasks
 - ▶ Different programs run on different cores and machines.
- ▶ For both paradigms, we rely on communication mechanisms to pass data around and to coordinate tasks (synchronization).

Shared Memory and Message Passing

- ▶ Communication via shared memory.
 - ▶ Data are shared in a common memory space that all cores are able to access.
 - ▶ Synchronization is the critical issue – need to use synchronization primitives like mutex and barriers.
- ▶ Communication via message passing.
 - ▶ Tasks communicate by sending and receiving messages only.
 - ▶ Synchronization is implied by the message.
 - ▶ Data must be received first before computation may start.
- ▶ Which one is more intuitive to work with in software? Which one is more close to actual hardware?

- ▶ Thread: what runs on a single core.
 - ▶ A program utilizes multiple cores by creating multiple threads.
 - ▶ Threads created from the same program may communicate through shared memory.
 - ▶ Multiple threads may run on a single core with the help of operating system, though for best computational efficiency, almost always a single thread is used per core.
- ▶ OS APIs to manage threads and to support synchronization are usually too cumbersome for parallel applications.
- ▶ OpenMP: An API for Writing Multithreaded Applications
 - ▶ A set of compiler directives and library routines for parallel application programmers
 - ▶ Greatly simplifies writing multi-threaded (MT) programs in Fortran, C and C++
 - ▶ Standardizes established SMP practice + vectorization and heterogeneous device programming

Outline

OpenMP Overview

Hello World!

Single Program Multiple Data

OpenMP Hello World!

```
// hw.cpp
#include <omp.h>
#include <stdio.h>

int main()
{
    #pragma omp parallel
    {
        printf("Hello, world!\n");
    }
    return 0;
}
```

- ▶ Compile: `g++ -fopenmp hw.cpp -o hw`
- ▶ Run: `./hw`
- ▶ Control number of threads: `OMP_NUM_THREADS=16 ./hw`

Know Who You Are

```
// hw2.cpp
#include <omp.h>
#include <stdio.h>

int main()
{
    #pragma omp parallel
    {
        int n = omp_get_num_threads();
        int i = omp_get_thread_num();
        printf("Thread %d of %d: Hello, world!\n", i, n);
    }
    return 0;
}
```

- ▶ Within an `omp parallel` block,
 - ▶ `omp_get_num_threads()` tells how many threads are there.
 - ▶ `omp_get_thread_num()` tells the index of this thread.

Anomaly

```
// hw3.cpp
#include <omp.h>
#include <stdio.h>

int main()
{
    #pragma omp parallel
    {
        int n = omp_get_num_threads();
        int i = omp_get_thread_num();
        printf("Thread %d of %d: ", i, n);
        printf("Hello, world!\n");
    }
    return 0;
}
```

- ▶ What output do you expect and what output do you see?
- ▶ The need for synchronization
 - ▶ Synchronization inside `printf()` ensures each call is completed atomically as a single operation.
 - ▶ However, different calls from different threads may interleave.
 - ▶ What about `std::cout`?

Fork-Join

```
// hw4.cpp
#include <omp.h>
#include <stdio.h>
int main() {
    #pragma omp parallel
    {
        int n = omp_get_num_threads(); int i = omp_get_thread_num();
        printf("Thread %d of %d: working on task 1\n", i, n);
    }
    printf("Done with task 1\n");
    #pragma omp parallel
    {
        int n = omp_get_num_threads(); int i = omp_get_thread_num();
        printf("Thread %d of %d: working on task 2\n", i, n);
    }
    printf("Done with task 2\n");
    return 0;
}
```

- There may be multiple `omp parallel` blocks and you may have usual sequential code in-between.
 - Only threads inside an `omp parallel` block are executed parallelly, and one must wait for all of them to finish before exiting the block.

Outline

OpenMP Overview

Hello World!

Single Program Multiple Data

Calculating π

```
// pi.cpp
#include <omp.h>
#include <stdio.h>
int main()
{
    const size_t num_steps= 4000000000LL;
    const double step = 1.0/num_steps;
    double sec = omp_get_wtime();
    double sum = 0;
    for (size_t i = 0; i < num_steps; ++i)
    {
        double x = (i+0.5)*step;
        sum = sum+4.0/(1.0+x*x);
    }
    double pi = step*sum;
    sec = omp_get_wtime()-sec;
    printf("pi = %.16f, time %.3f\n", pi, sec);
    return 0;
}
```

- ▶ Compile and optimize: `g++ -fopenmp -O2 pi.cpp -o pi`
- ▶ Use `omp_get_wtime()` to get the wall-clock time in seconds.

Calculating π with OpenMP

```
// pi_omp.cpp
...
    const int max_threads = 100;
    double sum[max_threads];
    int num_threads = 0;
    #pragma omp parallel
    {
        int n = std::min(omp_get_num_threads(), max_threads);
        int k = omp_get_thread_num();
        if (k < max_threads) {
            sum[k] = 0;
            for (size_t i = k; i < num_steps; i += n) {
                double x = (i+0.5)*step;
                sum[k] = sum[k]+4.0/(1.0+x*x);
            }
        }
        if (k == 0) num_threads = n;
    }

    double pi = 0;
    for (size_t k = 0; k < num_threads; ++k)
        pi += step*sum[k];
...
```

Synchronization

```
// pi_syn.cpp
...
double pi = 0;
#pragma omp parallel
{
    int n = omp_get_num_threads();
    int k = omp_get_thread_num();
    double sum = 0;
    for (size_t i = k; i < num_steps; i += n) {
        double x = (i+0.5)*step;
        sum = sum+4.0/(1.0+x*x);
    }
    #pragma omp critical
    {
        pi += step*sum;
    }
}
...
```

- ▶ We may simplify the previous code by using a local `sum`.
 - ▶ Update the shared `pi` from within the threads.
 - ▶ Use the `omp critical` block to ensure only one thread is allowed to update `pi` at a time.

Loop Parallelization

```
// pi_loop.cpp
...
double sum = 0;
#pragma omp parallel for reduction(+: sum)
for (size_t i = 0; i < num_steps; ++i)
{
    double x = (i+0.5)*step;
    sum = sum+4.0/(1.0+x*x);
}
double pi = step*sum;
...
```

- ▶ That loop is such a typical one that OpenMP can parallelize it automatically if you give a little bit of hint.
 - ▶ Use `for` in `omp parallel` to request for parallelization.
 - ▶ Use `reduction(+: sum)` to tell the compiler `sum` will be generated from the loop as a summation, and appropriate synchronization should be applied.

Summary and Advice

- ▶ OpenMP are widely available and widely used now.
 - ▶ Try different settings of `OMP_NUM_THREADS` with our Project 4–6 to see how `numpy` perform differently.
- ▶ Many embarrassingly parallel tasks like simple for loops can be parallelized by OpenMP automatically with a little bit of hint.
 - ▶ E.g. many tensor operators for machine learning.
 - ▶ The challenge is to optimize memory and cache performance.
- ▶ There is a function `omp_set_num_threads()` allowing to set how many threads you would like to use.
 - ▶ This is seldomly a good idea as consumers have very diverse choices of machines.
 - ▶ Your code should work on most settings and you should allow consumers to choose via `OMP_NUM_THREADS`.